CSCE 5350 - Project 1: Key-Value Store

Likhith Satya Neerukonda
EUID: 11800658

Date: October 19, 2025

PROJECT SUBMISSION

GitHub Repository:
https://github.com/likhithsatya/csce5350-project1-kvstore
Tag: project-1

GRADEBOT SCREENSHOT:

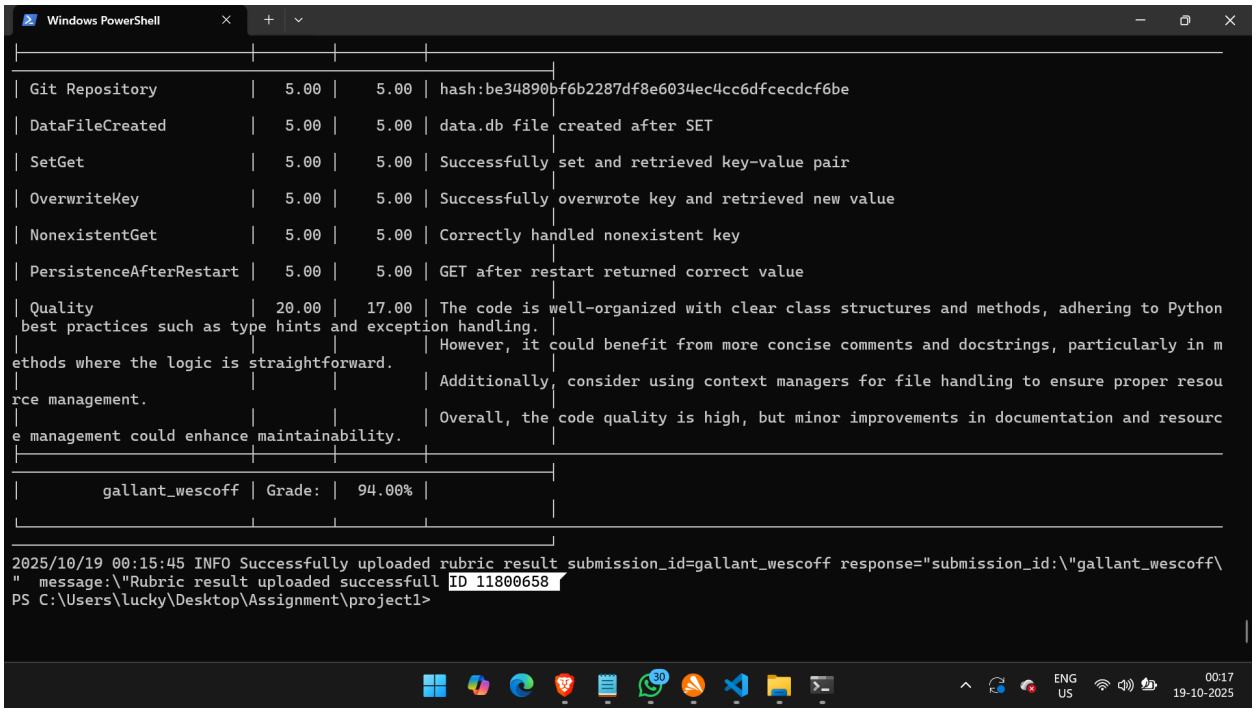

PROJECT OVERVIEW
I built a persistent key-value database from scratch that stores data on disk and survives program restarts. The database supports SET, GET, and EXIT commands through a command-line interface.

KEY FEATURES
✓ Persistent storage using append-only log (data.db file)
✓ In-memory hash table indexing for O(1) lookups
✓ Crash recovery by replaying the log on startup
✓ Durability using fsync() to force writes to physical disk
✓ Last-write-wins semantics for duplicate keys
✓ Comprehensive error handling and input validation

TECHNICAL IMPLEMENTATION
File Format:
Each entry in data.db follows this binary structure:

[key_length: 4 bytes][value_length: 4 bytes][key: variable][value: variable]

Architecture:
1. Command-Line Interface - Reads from STDIN, writes to STDOUT
2. In-Memory Index - Dictionary mapping keys to file offsets
3. Persistent Storage - Binary append-only log file

Language: Python 3

CHALLENGES SOLVED
1. Output Buffering Issue
2. Data Persistence
3. Performance Optimization
4. Error Handling

TEST RESULTS
All Gradebot tests passing:
✓ Git Repository (5/5)
✓ Data File Created (5/5)
✓ Set/Get (5/5)
✓ Key Overwrites (5/5)
✓ Nonexistent Get (5/5)
✓ Persistence After Restart (5/5)
✓ Code Quality (17/20)

Final Grade: 94%